



## Dependable Distributed Systems (9 CFU)


24/01/2023

### Exam A

Family Name \_\_\_\_\_ Name \_\_\_\_\_ Student ID \_\_\_\_\_

**Ex 1:** Provide the specification of the perfect failure detector and explain how it can be implemented in a synchronous system. In addition, discuss what does it happen to the correctness of the oracle if we use fair-loss point to point links instead of perfect ones.

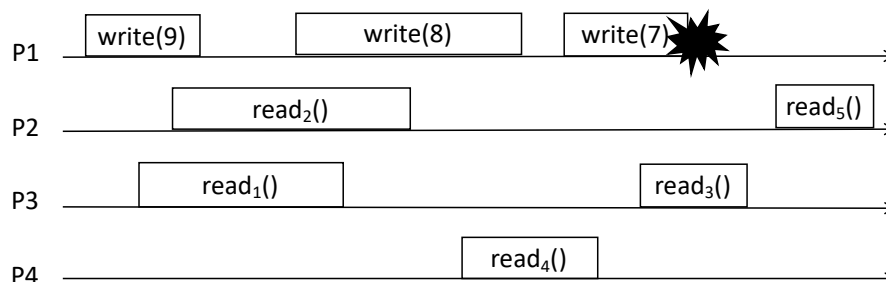
**Ex 2:** Let us consider the following execution history



Assess the truthfulness of every statement and provide a motivation for your answer:

1	The proposed run satisfies the Best Effort Broadcast specification	T	F
2	The strongest specification satisfied by the proposed run is Uniform Reliable Broadcast	T	F
3	If $p_3$ delivers $m_4$ , then the resulting run satisfies the Regular Reliable Broadcast specification	T	F
4	If $p_4$ delivers $m_4$ , then the resulting run satisfies the Regular Reliable Broadcast specification	T	F
5	If $p_2$ delivers $m_1$ , then the resulting run satisfies the Uniform Reliable Broadcast specification	T	F
6	The run satisfies $TO(NUA, WNUTO)$	T	F
7	The run satisfies the FIFO Broadcast specification	T	F
8	Let us assume $p_4$ does not deliver $m_2$ , then the resulting run satisfies $TO(NUA, WUTO)$	T	F
9	Let us assume $p_2$ crashes, then the resulting run satisfies $TO(UA, WNUTO)$	T	F
10	If $p_1$ does not deliver $m_1$ , then the resulting run satisfies $TO(UA, WUTO)$	T	F

**Ex 3:** Consider the execution depicted in the following figure



Answer the following questions:

- 1) Define ALL the values that can be returned by read operations (Rx) assuming the run refers to a regular register.
- 2) Define ALL the values that can be returned by read operations (Rx) assuming the run refers to an atomic register.

**Ex 4:** Let us assume a system working under a single workload component, the inter-arrival times and the service demands of such component are exponentially distributed, and the mean inter-arrival time is equal to 0.2 seconds.

- 1) Is it possible to estimate the expected response time if the mean service time is 0.05 sec? If yes, what is its value? If no, what is the minimum service time required to provide an estimation?
- 2) If the service is managed by two components that sequentially handle the requests, having service rates equal to  $\mu_1 = 8$  requests/sec and  $\mu_2 = 12$  requests/sec, is it possible to estimate the expected response time? If yes, what is its value? If no, explain why.
- 3) If every request must visit each component 2 times, namely, to be handled by the component twice, is it possible to estimate the expected response time? If yes, what is its value? If no, explain why.
- 4) Consider the setting presented at point 2. If you can replicate one component, splitting its load among two replicas, which component would you replicate to minimize the response time? What will be the expected response time?
- 5) Consider the configuration you provided at point 4 and assume that the MTBF of component 1 is 10 days, the MTBF of component 2 is 5 days and the repair rates of all components is 0.5 repair/day. Which setting guarantees higher availability between the setting at point 2 and the one you provided at point 4? Provide a motivation to your answer.

**NOTE:** It is sufficient to set the calculus to answer the questions.

**Ex 5:** Consider a distributed system composed of  $n$  processes. Processes are arranged in a ring topology. Every link of the ring is implemented through a perfect point-to-point channel. Each process in the system stores, in a local variable called `next`, the local identifier of the next process in the ring and maintains locally an integer value stored in a local variable `myVal`.

Let us consider the following specification:

### Module

**Name:** sum Oracle

### Events:

**Request:**  $\langle O_{\text{sum}}, \text{get\_sum} \rangle$ : invoke the oracle to compute the sum of values stored by correct processes

**Indication:**  $\langle O_{\text{sum}}, \text{sum\_return} \mid v \rangle$ : returns to the calling process the computed sum

### Properties:

- **Termination:** If a correct process invoke a `get_sum()` operation it eventually returns from the operation.
- **Validity:** If a correct process  $p_i$  returns a value  $v$ , then  $v$  is the sum of values stored by correct processes.

Answer to the following questions:

1. Assuming that (i) processes do not fail, (ii) each process is univocally identified by a unique integer identifier, (iii) the system is asynchronous and (iv) every process  $p_i$  does not know the overall number of processes in the system, write the pseudo-code of a distributed algorithm implementing the sum Oracle.

2. Assuming that (i) processes do not fail, (ii) processes are anonymous (i.e., every process has the same integer identifier), (iii) the system is asynchronous and (iv) every process  $p_i$  does not know the overall number of processes in the system, discuss if it is possible to implement the sum Oracle. If so, describe shortly the algorithm (no pseudo-code is needed), otherwise provide the intuition of the impossibility.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

Signature: \_\_\_\_\_

**Ex 1:** Provide the specification of the perfect failure detector and explain how it can be implemented in a synchronous system. In addition, discuss what does it happen to the correctness of the oracle if we use fair-loss point to point links instead of perfect ones.

**A FD IS A SW MODULE THAT PROVIDES INFORMATION ABOUT THE FAILURE STATUS OF PROCESSES. A PERFECT FD SATISFIES TWO PROPERTIES:**

**STRONG ACCURACY: CORRECT PROCESSES ARE NEVER ERRONEOUSLY SUSPECTED**  
**STRONG COMPLETENESS. EVERY PROCESS THAT CRASHES IS PERMANENTLY SUSPECTED BY EVERY CORRECT PROCESSES.**

**THESE PROPERTIES ENSURE THAT EACH FAILURE IS ALWAYS DETECTED AND THAT THERE ARE NO FALSE POSITIVES.**

**IN A SYNCHRONOUS SYSTEM, THE PFD CAN BE IMPLEMENTED BY EXPLOITING THE FACT THAT THERE ARE KNOWN LIMITS ON THE COMMUNICATION AND PROCESSING TIME BETWEEN THE PROCESSES THIS MEANS THAT, FOR EACH PROCESS, IT'S POSSIBLE TO ESTABLISH A MAXIMUM INTERVAL WITHIN WHICH AN ANSWER MUST BE RECEIVED FROM ANOTHER PROCESS TO CONSIDER IT STILL ACTIVE.**


**THIS IMPLEMENTATION IS BASED ON A HEARTBEAT AND TIMEOUT MECHANISM: EACH PROCESS PERIODICALLY SENDS MESSAGE OF "VITALITY" (HRBT) TO OTHERS AND, IF IT DOESN'T RECEIVE A RESPONSE WITHIN A MAXIMUM PREDEFINED TIME, CAN CONCLUDE WITH CERTAINTY THAT THE REMOTE PROCESS IS FAULTY.**

**USING FAIR-LOSS P2PL, MESSAGES CAN BE OCCASIONALLY LOST:**

**STRONG ACCURACY PROBLEM: IF A PROCESS IS CORRECT BUT ITS HRBT IS LOST, ANOTHER PROCESS COULD ERRONEOUSLY SUSPECT IT TO BE FAULTY.**

**STRONG COMPLETENESS: IF P CRASHES AND THE MSG OF FAILURE LOST, HOWEVER P WILL NO LONGER RESPOND WITH HEARTBEAT REPLY.**

Ex 2: Let us consider the following execution history

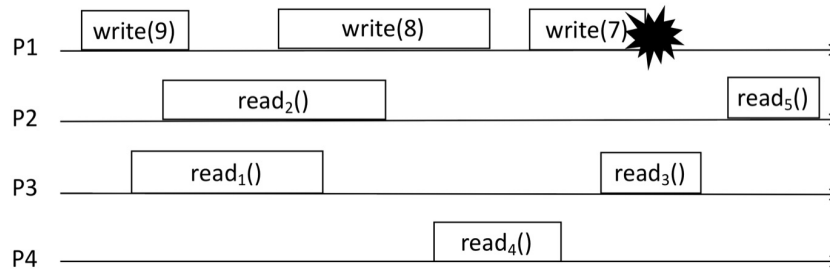


Assess the truthfulness of every statement and provide a motivation for your answer:

1	The proposed run satisfies the Best Effort Broadcast specification	<input checked="" type="checkbox"/>	F
2	The strongest specification satisfied by the proposed run is Uniform Reliable Broadcast	T	<input checked="" type="checkbox"/>
3	If $p_3$ delivers $m_4$ , then the resulting run satisfies the Regular Reliable Broadcast specification	<input checked="" type="checkbox"/>	F
4	If $p_4$ delivers $m_4$ , then the resulting run satisfies the Regular Reliable Broadcast specification	T	<input checked="" type="checkbox"/>
5	If $p_2$ delivers $m_1$ , then the resulting run satisfies the Uniform Reliable Broadcast specification	T	<input checked="" type="checkbox"/>
6	The run satisfies TO(NUA, WNUTO)	T	<input checked="" type="checkbox"/>
7	The run satisfies the FIFO Broadcast specification	T	<input checked="" type="checkbox"/>
8	Let us assume $p_4$ does not deliver $m_2$ , then the resulting run satisfies TO(NUA, WUTO)	T	<input checked="" type="checkbox"/>
9	Let us assume $p_2$ crashes, then the resulting run satisfies TO(UA, WNUTO)	T	<input checked="" type="checkbox"/>
10	If $p_1$ does not deliver $m_1$ , then the resulting run satisfies TO(UA, WUTO)	T	<input checked="" type="checkbox"/>

- 1)  $p_2$  AND  $p_4$  SATISFY VALIDITY, NO DUPLICATION AND NO CREATION.
- 2)  $p_1$  (FAULTY) DELIVERS  $m_1$ , WHICH IS NOT DELIVERED BY  $p_2$  AND  $p_4$  (CORRECT).
- 3)  $p_2$  AND  $p_4$  (CORRECT) SATISFIE AGREEMENT.
- 4)  $p_2$  DOESN'T DELIVER  $m_4$  (NO AGREEMENT).
- 5)  $p_4$  DOESN'T DELIVER  $m_1$  (NO UNIFORM AGREEMENT).
- 6) IT'S NOT TO BECAUSE  $p_2$  AND  $p_4$  HAVE DIFFERENT ORDERS.
- 7)  $m_2 \rightarrow m_3$  NOT SATISFIED BY  $p_2$
- 8)  $p_2$  AND  $p_4$  HAVE DIFFERENT SET OF DELIVERY
- 9)  $p_4$  DOESN'T DELIVER  $m_1$  (NO UNIFORM AGREEMENT).
- 10) IT'S NOT TO BECAUSE  $p_2$  AND  $p_4$  HAVE DIFFERENT ORDERS.

Ex 3: Consider the execution depicted in the following figure



Answer the following questions:

- 1) Define ALL the values that can be returned by read operations (Rx) assuming the run refers to a regular register.
- 2) Define ALL the values that can be returned by read operations (Rx) assuming the run refers to an atomic register.

1)  $R_1(): 0, 9, 8$   
 $R_2(): 0, 9, 8$   
 $R_3(): 8, 7$   
 $R_4(): 9, 8, 7$   
 $R_5(): 8, 7$

2)  $R_1(): 0, 9, 8$   
 $R_2(): 0, 9, 8$   
 $R_4(): \begin{cases} 9, 8, 7 & \text{IF } R_2(): 0, 9 \text{ AND } R_1(): 0, 9 \\ 8, 7 & \text{IF } R_2(): 8 \text{ AND } R_1(): 8 \end{cases}$   
 $R_3(): \begin{cases} 8, 7 & \text{IF } R_4(): 9, 8 \\ 7 & \text{IF } R_4(): 7 \end{cases}$   
 $R_5(): \begin{cases} 8, 7 & \text{IF } R_3(): 8 \\ 7 & \text{IF } R_3(): 7 \end{cases}$

**Ex 5:** Consider a distributed system composed of  $n$  processes. Processes are arranged in a ring topology. Every link of the ring is implemented through a perfect point-to-point channel. Each process in the system stores, in a local variable called `next`, the local identifier of the next process in the ring and maintains locally an integer value stored in a local variable `myVal`.

Let us consider the following specification:

#### Module

**Name:** sum Oracle

#### Events:

**Request:**  $\langle O_{sum}, get\_sum \rangle$ : invoke the oracle to compute the sum of values stored by correct processes



**Indication:**  $\langle O_{sum}, sum\_return \mid v \rangle$ : returns to the calling process the computed sum

#### Properties:

- **Termination:** If a correct process invoke a `get_sum()` operation it eventually returns from the operation.
- **Validity:** If a correct process  $p_i$  returns a value  $v$ , then  $v$  is the sum of values stored by correct processes.

Answer to the following questions:

1. Assuming that (i) processes do not fail, (ii) each process is univocally identified by a unique integer identifier, (iii) the system is asynchronous and (iv) every process  $p_i$  does not know the overall number of processes in the system, write the pseudo-code of a distributed algorithm implementing the sum Oracle.
2. Assuming that (i) processes do not fail, (ii) processes are anonymous (i.e., every process has the same integer identifier), (iii) the system is asynchronous and (iv) every process  $p_i$  does not know the overall number of processes in the system, discuss if it is possible to implement the sum Oracle. If so, describe shortly the algorithm (no pseudo-code is needed), otherwise provide the intuition of the impossibility.



1) PROCESS DON'T FAIL, EACH  $p$  HAS AN ID, SYSTEM ASYNCL,  $N$  UNKNOWN

UPON EVENT  $\langle SO, INIT \rangle$  DO

    NEXT = GET.NEXT()

    MYVAL

UPON EVENT  $\langle SO, GET.SUM() \rangle$  DO

    TRIGGER  $\langle PP2PL, SEND \mid [MYVAL, SELF] \rangle$  TO NEXT

UPON EVENT  $\langle PP2PL, DELIVER \mid [VAL, p] \rangle$  DO

    IF SELF ==  $p$  THEN **SIAMO TORNATI AL PRIMO**

        TRIGGER  $\langle SO, SUM.RETURN \mid VAL \rangle$

    ELSE

        TRIGGER  $\langle PP2PL, SEND \mid [VAL + MYVAL, p] \rangle$  TO NEXT

2) IMPOSSIBLE BECAUSE IN THE DELIVERY WE CHECK IF ALL PROCESSES SEND THEIR VALUE BY USING THEIR IDS. WITHOUT ID WE DON'T KNOW WHEN TO STOP.