



Dependable Distributed Systems (9 CFU)

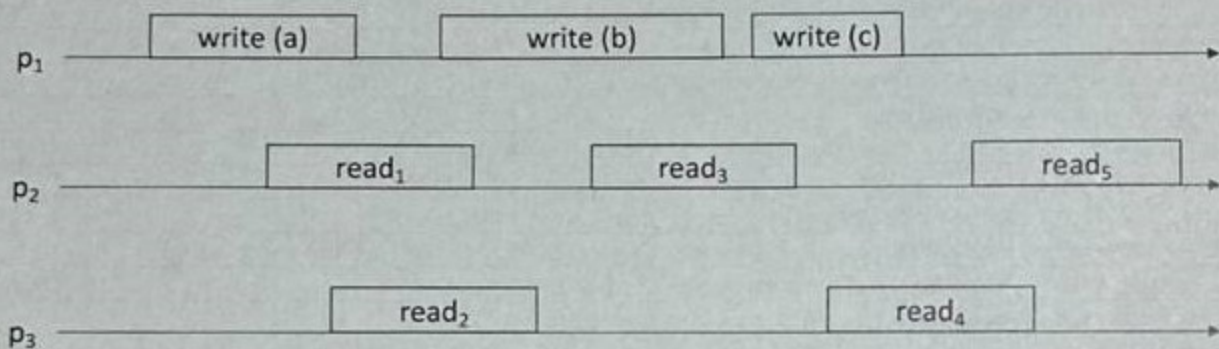
14/02/2023

Exam A

Family Name _____ Name _____ Student ID _____

Ex 1: Provide the specification of the regular consensus primitive and describe the implementation presented during the lectures in synchronous systems. Finally, discuss its performance (in terms of number of messages exchanged) to reach consensus.

Ex 2: Let us consider the following execution history




Assuming that the initial value stored in the register is 0, assess the truthfulness of every statement and provide a motivation for your answer:

1	If the proposed run refers to a regular register, then 0 is a valid value for $read_1$	T	F
2	If the proposed run refers to a regular register, then 0 is not a valid value for $read_2$	T	F
3	If the proposed run refers to a regular register, then a is a valid value for $read_4$	T	F
4	If the proposed run refers to a regular register, then b is a valid value all read operations	T	F
5	If the proposed run refers to a regular register, then $read_5$ may return only b and c	T	F
6	If the proposed run refers to an atomic register, then $read_1$ and $read_2$ must return the same value	T	F
7	If the proposed run refers to an atomic register, then $read_3$ returns b if and only if $read_1$ returns b	T	F
8	If the proposed run refers to an atomic register, then $read_4$ and $read_5$ must always return the same value	T	F
9	If the proposed run refers to an atomic register and $read_4$ returns c , then $read_3$ necessarily returned c	T	F
10	If the proposed run refers to an atomic register and $read_1$ returned b then $read_2$ can return only the value b	T	F

Ex 3: Consider the broadcast communication primitives studied during the course, assess the truthfulness of every statement below and for each statement provide a motivation (also using suitable examples).

1	If a run R satisfies Causal Order Broadcast, then R satisfies FIFO order Broadcast	T	F
2	If a run R satisfies Total Order Broadcast, then R satisfies Regular Reliable Broadcast	T	F
3	It is not possible to define a run R that satisfies both Total Order Broadcast and FIFO order Broadcast	T	F
4	If a run R satisfies TO (NUA, WUTO), then R satisfies also Uniform Reliable Broadcast	T	F
5	If a run R satisfies TO (NUA, WUTO), then R satisfies TO (NUA, WNUTO)	T	F

Ex 4: Let us consider the following execution history



Let us denote with $ck(e_i)$ the logical clock associated to event e_i . Considering the execution history shown in the figure above, assess the truthfulness of every statement and provide a motivation for your answer:

1	If we use scalar clocks for timestamping events, then $ck(e_6) > ck(e_5)$	T	F
2	e_2 happened before e_3 (according with Lamport's definition of happened-before)	T	F
3	If we use scalar clocks for timestamping events, then $ck(e_{14}) = ck(e_{10}) + 1$	T	F
4	e_6 and e_7 are concurrent events	T	F
5	If we use scalar clocks for timestamping events, then $ck(e_9) < ck(e_{11})$	T	F
6	If we use vector clocks for timestamping events, then $ck(e_{13}) = [4, 6, 4]$	T	F
7	If we use vector clocks for timestamping events, then $ck(e_1) = ck(e_3)$	T	F
8	If we use vector clocks for timestamping events, then $ck(e_5)$ and $ck(e_{10})$ are not comparable	T	F
9	If we use vector clocks for timestamping events, then $ck(e_8) < ck(e_{10})$	T	F
10	If we use vector clocks for timestamping events, then $ck(e_4) = [1, 1, 1]$	T	F

Ex 5: Let us consider a distributed system composed of a set $C = \{c_1, c_2, \dots, c_m\}$ of clients and a set $R = \{r_1, r_2, \dots, r_n\}$ of replicas. Clients and replicas are univocally identified by an integer. Replicas communicate among themselves by message passing and are arranged in a unidirectional ring topology. Every replica r_i can send messages (over a perfect point-to-point link) only to its neighbor whose name is stored in a local variable $next_i$. Replicas may fail by crash and every replica r_i is equipped with a perfect oracle that notifies, through the $new_next(r_j)$ event, the identifier r_j of the new r_i 's neighbor in case of failure.

Replicas need to maintain a shared object by implementing the primary-backup replication schema.

Every client may communicate with replicas using perfect point-to-point links. Initially, clients only know the identifier of the current primary and they store it in a local variable $primary$. So, when they need to issue operations, they just need to interact with it.

Given the current scenario, answer the following points:

1. Write the pseudo-code of the distributed protocol implementing the replication schema (both client and replicas side).
2. Assuming a long enough period $[t, t']$ without any failure, compute the expected response time (from the client point of view) to execute an operation op knowing that the inter arrival time between requests is exponentially distributed with parameter $\lambda = 2$ req/sec, that the average service time to execute op on a replica is $1/4$ sec, that all P2P links have a service time of $1/7$ sec per message, and that all service times are exponentially distributed (the upstream and downstream of a P2P link can be assumed independent).
3. Assuming that any replica fails every 24h on average and that the replicas take around 30 minutes to elect a new primary, provide an estimation of the steady-state availability of the shared object considering the previous times exponentially distributed.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

Signature: _____

Ex 1: Provide the specification of the regular consensus primitive and describe the implementation presented during the lectures in synchronous systems. Finally, discuss its performance (in terms of number of messages exchanged) to reach consensus.

IT'S A MECHANISM THAT ALLOWS A SET OF PROCESSES IN A DISTRIBUTED SYSTEM TO AGREE ON A SINGLE VALUE, EVEN IN PRESENCE OF FAILURES. IT IS BASED ON:

VALIDITY. IF A PROCESS DECIDES v , THEN v WAS PROPOSED BY SOME PROCESSES.

INTEGRITY: NO PROCESS DECIDES TWICE.

AGREEMENT: NO TWO CORRECT PROCESSES DECIDE DIFFERENTLY.

TERMINATION: EVERY CORRECT PROCESS EVENTUALLY DECIDES SOME VALUE.


IT'S IMPOSSIBLE TO GUARANTEE CONSENSUS IN AN ASYNCHRONOUS SYSTEM

FLOODING CONSENSUS IS AN ALGORITHM BASED ON A PROPAGATION FOR FLOODING MESSAGES TO ENSURE THAT ALL CORRECT PROCESSES REACH AN AGREEMENT ON THE VALUE TO BE DECIDED.

EACH PROCESS BEGINS BY PROPOSING A VALUE AND TRANSMITS IT TO OTHERS USING beb . THE PROCESSES COLLECT THE PROPOSALS RECEIVED IN THE FOLLOWING ROUNDS, UPDATING A SET OF CANDIDATE VALUES. IF A PROCESS RECEIVES MESSAGES FROM ALL THE NODES FROM WHICH HE HAD RECEIVED IN THE PREVIOUS ROUND, HE CHOOSES THE MINIMUM VALUE BETWEEN THOSE PROPOSED AND PROPAGATES IT AS A FINAL DECISION. THE PROCESS ENDS WHEN A DECISION MESSAGE IS DECIDED, WHICH IS IMMEDIATELY PROPAGATED TO GUARANTEE CONVERGENCE.

THE ALGORITHM HAS A COMPLEXITY EQUAL TO $O(n^2)$, WITH n # OF PROCESSES. EACH OF THEM SENDS ITS VALUE TO ALL THE OTHERS, AND EACH NODE FORWARDS THE VALUES RECEIVED, CAUSING A FLOODING EFFECT.

Ex 2: Let us consider the following execution history



Assuming that the initial value stored in the register is 0, assess the truthfulness of every statement and provide a motivation for your answer:

1	If the proposed run refers to a regular register, then 0 is a valid value for read ₁	<input checked="" type="checkbox"/>	F
2	If the proposed run refers to a regular register, then 0 is not a valid value for read ₂	<input type="checkbox"/>	T
3	If the proposed run refers to a regular register, then <i>a</i> is a valid value for read ₄	<input type="checkbox"/>	T
4	If the proposed run refers to a regular register, then <i>b</i> is a valid value all read operations	<input type="checkbox"/>	T
5	If the proposed run refers to a regular register, then read ₃ may return only <i>b</i> and <i>c</i>	<input type="checkbox"/>	T
6	If the proposed run refers to an atomic register, then read ₁ and read ₂ must return the same value	<input type="checkbox"/>	T
7	If the proposed run refers to an atomic register, then read ₃ returns <i>b</i> if and only if read ₁ returns <i>b</i>	<input type="checkbox"/>	T
8	If the proposed run refers to an atomic register, then read ₄ and read ₅ must always return the same value	<input type="checkbox"/>	T
9	If the proposed run refers to an atomic register and read ₄ returns <i>c</i> , then read ₃ necessarily returned <i>c</i>	<input type="checkbox"/>	T
10	If the proposed run refers to an atomic register and read ₁ returned <i>b</i> then read ₂ can return only the value <i>b</i>	<input checked="" type="checkbox"/>	F


- 1) $R_1(): 0, a, b$
- 2) $R_2(): 0, a, b$
- 3) $R_4(): b, c$
- 4) $R_5(): c$
- 5) $R_3(): a, b, c$
- 6) R_1 AND R_2 ARE CONCURRENT
- 7) $R_3(): b$ ALSO IF $R_1(): a$
- 8) R_4 CAN RETURN b BUT R_5 ALWAYS c
- 9) R_3 CAN RETURN a OR b
- 10) R_1 AND R_2 ARE CONCURRENT, R_2 IS THE SUBSEQUENT OF R_1 AND CANNOT READ c

Ex 3: Consider the broadcast communication primitives studied during the course, assess the truthfulness of every statement below and for each statement provide a motivation (also using suitable examples).

1	If a run R satisfies Causal Order Broadcast, then R satisfies FIFO order Broadcast	X	F
2	If a run R satisfies Total Order Broadcast, then R satisfies Regular Reliable Broadcast	X	F
3	It is not possible to define a run R that satisfies both Total Order Broadcast and FIFO order Broadcast	T	X
4	If a run R satisfies TO (NUA, WUTO), then R satisfies also Uniform Reliable Broadcast	T	X
5	If a run R satisfies TO (NUA, WUTO), then R satisfies TO (NUA, WNUTO)	X	F

1) CASUAL ORDER = FIFO ORDER + LOCAL ORDER


2) TO IS A RB ORDER WITH ORDER PROPERTIES



4) NUA: FAULTY PROCESSES DELIVER MSGs THAT ARE NOT DELIVERED BY THE CORRECT.
UNIFORM AGREEMENT: IF A PROCESS (CORRECT OR FAULTY) DELIVER m , EVERY CORRECT PROCESS DELIVERS m



5) WNUTO IS A SUBSET OF WUTO

Ex 4: Let us consider the following execution history



Let us denote with $ck(e_i)$ the logical clock associated to event e_i . Considering the execution history shown in the figure above, assess the truthfulness of every statement and provide a motivation for your answer:


1	If we use scalar clocks for timestamping events, then $ck(e_6) > ck(e_5)$	T	X
2	e_2 happened before e_3 (according with Lamport's definition of happened-before)	T	X
3	If we use scalar clocks for timestamping events, then $ck(e_{14}) = ck(e_{10}) + 1$	X	F
4	e_6 and e_7 are concurrent events	X	F
5	If we use scalar clocks for timestamping events, then $ck(e_9) < ck(e_{11})$	T	X
6	If we use vector clocks for timestamping events, then $ck(e_{13}) = [4, 6, 4]$	T	X
7	If we use vector clocks for timestamping events, then $ck(e_1) = ck(e_3)$	T	X
8	If we use vector clocks for timestamping events, then $ck(e_5)$ and $ck(e_{10})$ are not comparable	T	X
9	If we use vector clocks for timestamping events, then $ck(e_8) < ck(e_{10})$	T	X
10	If we use vector clocks for timestamping events, then $ck(e_4) = [1, 1, 1]$	T	X



Ex 5: Let us consider a distributed system composed of a set $C = \{c_1, c_2, \dots, c_m\}$ of clients and a set $R = \{r_1, r_2, \dots, r_n\}$ of replicas. Clients and replicas are univocally identified by an integer. Replicas communicate among themselves by message passing and are arranged in a unidirectional ring topology. Every replica r_i can send messages (over a perfect point-to-point link) only to its neighbor whose name is stored in a local variable $next_i$. Replicas may fail by crash and every replica r_i is equipped with a perfect oracle that notifies, through the $new_next(r_i)$ event, the identifier r_j of the new r_i 's neighbor in case of failure. Replicas need to maintain a shared object by implementing the primary-backup replication schema. Every client may communicate with replicas using perfect point-to-point links. Initially, clients only know the identifier of the current primary and they store it in a local variable $primary$. So, when they need to issue operations, they just need to interact with it.

Given the current scenario, answer the following points:

1. Write the pseudo-code of the distributed protocol implementing the replication schema (both client and replicas side).



REPLICAS:

UPON EVENT $\langle R, INIT \rangle$ DO

```

NEXT = SELF.NEXT()
VAL.x = DEFAULT
CORRECT =  $\perp$ 
PRIMARY = MAX_RANK(CORRECT)

```

UPON EVENT $\langle R, OPERATION \mid v \rangle$ DO

```

VAL.x = EXECUTE OP (v)
TRIGGER  $\langle PP2PL, SEND \mid [UPDATE, VAL.x, SELF] \rangle$  TO NEXT

```

UPON EVENT $\langle PP2PL, DELIVER \mid [UPDATE, VAL.x, p] \rangle$

```

IF p = SELF THEN
    TRIGGER  $\langle R, OK \mid OP \rangle$ 
ELSE
    VAL.x = VAL
    TRIGGER  $\langle PP2PL, SEND \mid [UPDATE, VAL.x, p] \rangle$ 

```

UPON EVENT $\langle PO, NEW.NEXT \mid NEXT \rangle$ DO

```

FAIL = NEXT
CORRECT = CORRECT  $\setminus \{FAIL\}$ 
IF FAIL == PRIMARY THEN
    PRIMARY = MAX_RANK(CORRECT)
ELSE:
    TRIGGER  $\langle PP2PL, SEND \mid [UPDATE, VAL.x, p] \rangle$ 
    TRIGGER  $\langle PP2PL, SEND \mid [CRASH, FAIL] \rangle$  TO NEXT

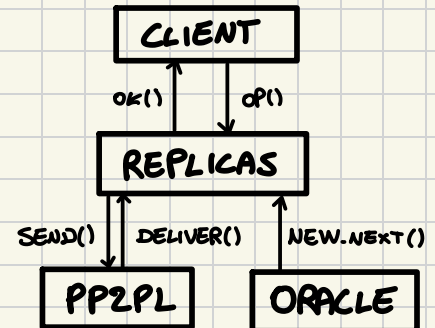
```

UPON EVENT $\langle PP2PL, DELIVER \mid [CRASH, FAIL] \rangle$

```

CORRECT = CORRECT'  $\setminus \{FAIL\}$ 
TRIGGER  $\langle PP2PL, SEND \mid [CRASH, FAIL] \rangle$  TO NEXT
IF FAIL == PRIMARY
    PRIMARY = MAX_RANK(CORRECT)

```



CLIENT:

UPON EVENT $\langle CL, INIT \rangle$ DO
OPERATIONS = \emptyset

UPON EVENT $\langle CL, DELIVER \mid REQUEST, OP \rangle$ DO
OPERATIONS = OPERATIONS $\cup \{OP\}$
IF $|OPERATIONS| = 1$ THEN
STARTTIMER (TIMER)

UPON EVENT $\langle CL, OK \mid OP \rangle$ DO
OPERATIONS = OPERATIONS $\setminus \{OP\}$

UPON EVENT $\langle TIMEOUT \rangle$
IF OPERATIONS $\neq \emptyset$ THEN
TRIGGER $\langle PP2PL, SEND \mid OP, OPERATION(i) \rangle$
STARTTIMER (TIMER)