


## 2025-01-29 (Exam A)

Ex. 1: Provide the specification of the leader election primitive, describe its implementation discussion in detail the system model

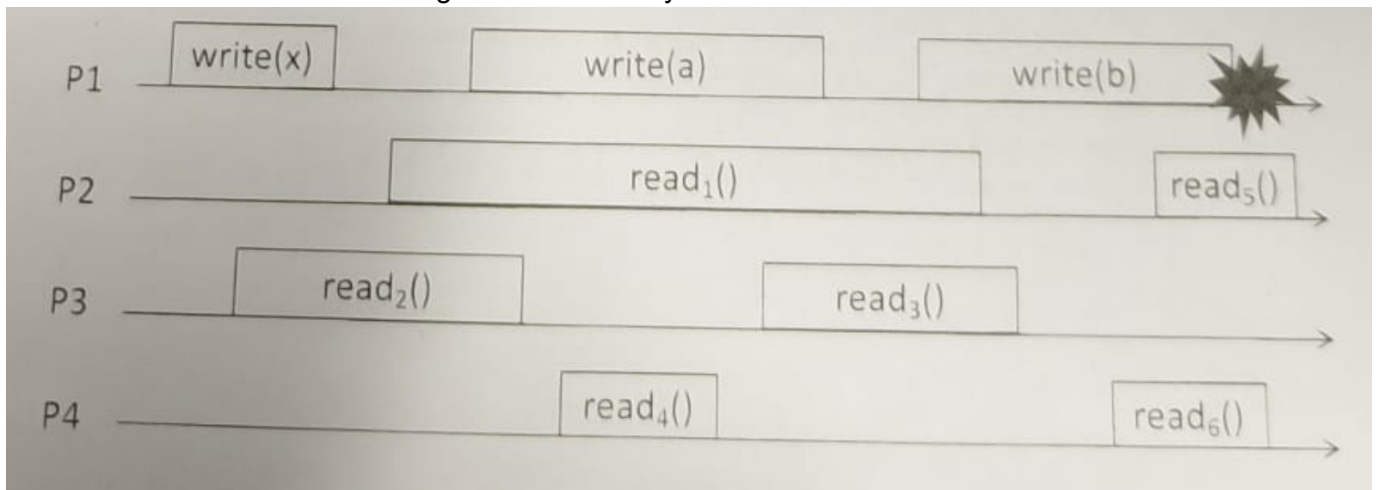
Ex 2. Let us consider the following execution history



Asses the truthfulness of every statement and provide a motivation for your answer:

1. The proposed run satisfies the Best Effort Broadcast specification
2. The strongest specification satisfied by the proposed run is Uniform Reliable Broadcast
3. If  $p_3$  delivers  $m_1$ , then the resulting run satisfies the Uniform Reliable Broadcast specification
4. If  $p_3$  does not deliver  $m_5$ , then the resulting run satisfies the Uniform Reliable Broadcast specification
5. If  $p_3$  does not crash, then the resulting run satisfies the Best Effort Broadcast specification
6. The run satisfies  $TO(NUA, WNUTO)$
7. The run satisfies the FIFO Broadcast specification
8. Let us assume  $p_3$  does not deliver  $m_5$ , then the resulting run satisfies  $TO(UA, WNUTO)$
9. Let us assume  $p_2$  does not crash, then the resulting run satisfies  $TO(NUA, WNUTO)$
10. If  $p_2$  does not deliver  $m_2$  and  $m_3$ , then the resulting run satisfies  $TO(NUA, WUTO)$

Ex. 3: Let us consider the following execution history



Assuming that the initial value stored in the register is 0, assess the truthfulness of every statement and provide a motivation for your answer:

1. If the proposed run refers to a regular register, then 0 is a valid value for  $read_1$
2. If the proposed run refers to a regular register, then 0 is not a valid value for  $read_2$
3. If the proposed run refers to a regular register, then a is a valid value for  $read_4$
4. If the proposed run refers to a regular register, then x can be returned only by  $read_2$

5. If the proposed run refers to a regular register, then  $\text{read}_5$  may return only b
6. If the proposed run refers to an atomic register, then  $\text{read}_3$  and  $\text{read}_2$  must return the same value
7. If the proposed run refers to an atomic register, then  $\text{read}_5$  returns b if and only if  $\text{read}_3$  returns b
8. If the proposed run refers to an atomic register, then  $\text{read}_6$  and  $\text{read}_5$  must always return the same value
9. If the proposed run refers to an atomic register and  $\text{read}_3$  returns b, then  $\text{read}_6$  necessarily returned b
10. If the proposed run refers to an atomic register and  $\text{read}_2$  returned a then  $\text{read}_4$  can return only the value a

Ex. 4: Let us consider the following algorithm implementing a publish-subscribe primitive

#### Algorithm 1: Publisher code

```

1 Init
2  $\text{brokers}[] \leftarrow \text{get\_brokers}();$  /* array of length  $|T|$  where in the  $i$ -th entry the
   publisher stores the set of brokers managing topic  $t_i$  */
3 upon event  $\text{publish}(e_i, t_i)$ 
4   foreach  $b_j \in \text{broker}[t_i]$  do
5     | trigger  $\text{pp2pSend}(\text{PUBLISH}, e_i, t_i)$  to  $b_j$ ;
6   end

```

#### Algorithm 2: Subscriber code

```

1 Init
2  $\text{brokers}[] \leftarrow \text{get\_brokers}();$  /* array of length  $|T|$  where in the  $i$ -th entry the
   subscriber stores the set of brokers managing topic  $t_i$  */
3  $\text{my\_subscription} \leftarrow \text{get\_my\_initial\_subscription}();$  /* set of topics */
4  $\text{notified}[] \leftarrow \{\emptyset\}^{|T|};$  /* array of length  $|T|$  where in the  $i$ -th entry the
   subscriber stores the set of notified events for topic  $t_i$  */
5 upon event  $\text{subscribe}(t_i)$ 
6    $\text{my\_subscription} \leftarrow \text{my\_subscription} \cup \{t_i\};$ 
7   foreach  $b_j \in \text{brokers}[t_i]$  do
8     | trigger  $\text{pp2pSend}(\text{SUBSCRIBE}, s_i, t_i)$  to  $b_j$ ;
9   end
10 upon event  $\text{unsubscribe}(t_i)$ 
11    $\text{my\_subscription} \leftarrow \text{my\_subscription} \setminus \{t_i\};$ 
12   foreach  $b_j \in \text{brokers}[t_i]$  do
13     | trigger  $\text{pp2pSend}(\text{UNSUBSCRIBE}, s_i, t_i)$  to  $b_j$ ;
14   end
15 upon event  $\text{pp2pDeliver}(\text{NOTIFY}, e_j, t_j)$ 
16 if  $t_j \in \text{my\_subscription}$  AND  $e_j \notin \text{notified}[t_j]$  then
17   | trigger  $\text{NOTIFY}(e_j, t_j)$ ;
18   |  $\text{notified} \leftarrow \text{notified} \cup \{(e_j, t_j)\};$ 
19 end

```

#### Algorithm 3: Broker code

```

1 Init
2  $\text{subscription}[] \leftarrow \{\emptyset\}^{|T|};$  /* array of length  $|T|$  where in the  $i$ -th entry the
   broker stores the set of subscribers for topic  $t_i$  */
3  $\text{topic\_to\_manage} \leftarrow \text{get\_topic\_to\_manage}();$  /* set of topics that the current
   broker must manage */
4 upon event  $\text{pp2pDeliver}(\text{PUBLISH}, e_i, t_i)$ 
5 if  $t_i \in \text{topic\_to\_manage}$  then
6   | foreach  $s_k \in \text{subscription}[t_i]$  do
7     | | trigger  $\text{pp2pSend}(\text{NOTIFY}, e_i, t_i)$  to  $s_k$ ;
8   | end
9 end
10 upon event  $\text{pp2pDeliver}(\text{SUBSCRIBE}, s_i, t_i)$ 
11    $\text{subscription}[t_i] \leftarrow \text{subscription}[t_i] \cup \{s_i\};$ 
12 upon event  $\text{pp2pDeliver}(\text{UNSUBSCRIBE}, s_i, t_i)$ 
13    $\text{subscription}[t_i] \leftarrow \text{subscription}[t_i] \setminus \{s_i\};$ 

```

Let us assume that

1. the underlying communication system is synchronous and that every message sent over perfect point to point link takes at most  $\delta$  time to be delivered.
2. links guarantee the FIFO property
3. for each topic  $t_i$ , there are at least 3 brokers managing it

Assess the truthfulness of every statement and provide a motivation for your answer:

1. If subscriptions are not modified and there is at most one failure in the system, then every event  $e$  published on a topic  $t_i$  will be eventually notified to all subscribers subscribed to it
2. If subscriptions are not modified and there are no failures in the system, then any pair of subscribers  $s_i$  and  $s_j$  both subscribed to the same topic  $t_i$  will deliver the same set of events published on  $t_i$
3. If subscriptions are not modified and there is at most one failure in the system, if a subscriber  $s_i$  notifies an event  $e$  before an event  $e'$  (both published on a topic  $t_i$ ), then any other subscriber having  $t_i$  in its subscriptions will notify  $e$  before  $e'$
4. In absence of failures, every event published on a topic  $t_i$  at time  $\tau$  will be eventually notified to every subscriber having  $t_i$  in its subscription at time  $\tau$
5. In absence of failures, if a subscriber  $s_i$  unsubscribes a topic  $t_i$  at time  $\tau$ , it will not notify any event for topic  $t_i$  after time  $\tau$

Ex. 5:

Let us consider a distributed system composed of  $n$  processes  $p_1, p_2, \dots, p_n$  partitioned in two disjoint sets acting as clients and replicas. Each process is identified by a unique integer identifier and can communicate with other processes by exchanging messages using perfect point-to-point links. Processes collaborate to implement a replication protocol emulating a linearizable append-only log.

The append-only log is characterized by the following operations:

- **Append( $v$ )**: appends the value  $v$  to the queue of the log
- **Get**: returns the content of the log, e.g. the list of values in the order in which they have been appended.

Assuming that:

1. processes may fail by crash
2. Every pair of processes can communicate exchanging messages over perfect point-to-point links
3. All processes have access to a perfect failure detector
4. Replicas have access to a uniform reliable broadcast primitive

Write the pseudocode of the distributed algorithm implementing the append-only log.


Ex. 1: Provide the specification of the leader election primitive, describe its implementation discussion in detail the system model

ELECTING A LEADER CONSISTS OF IDENTIFYING A PROCESS THAT ACTS AS A COORDINATOR BETWEEN DISTRIBUTED PROCESSES. A LEADER CAN MANAGE EXCLUSIVE ACCESS TO A SHARED RESOURCE.

IN A SYNCHRONOUS SYSTEM, A POSSIBLE APPROACH IS TO ASSIGN EACH PROCESS A UNIQUE ID AND ENSURE THAT THE PROCESS WITH THE HIGHEST ID BECOMES THE LEADER, ENSURING THAT ALL THE PROCESSES ELECT THE SAME LEADER IN A FINISHED NUMBER OF STEPS.

IN AN EVENTUALLY SYNCHRONOUS SYSTEM, THE ABSENCE OF A WELL KNOWN LIMIT ON COMMUNICATION TIMES AND THE DETECTION OF FAULTS PREVENTS THE CONSTRUCTION OF A PERFECT ORACLE. A PROCESS COULD BE ERRONEOUSLY CONSIDERED FAILED AND REPLACED AS A LEADER, LEADING TO INFINITE CHANGES IN THE ROLE OF LEADER WITHOUT EVER ACHIEVING A DEFINITE STABILIZATION.

Ex 2. Let us consider the following execution history

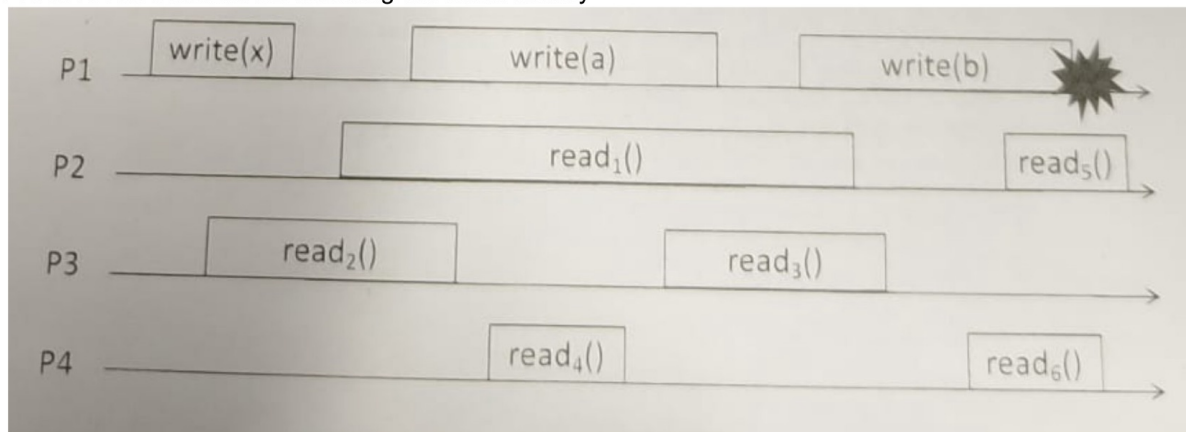


Assess the truthfulness of every statement and provide a motivation for your answer:

1. The proposed run satisfies the Best Effort Broadcast specification
2. The strongest specification satisfied by the proposed run is Uniform Reliable Broadcast
3. If  $p_3$  delivers  $m_1$ , then the resulting run satisfies the Uniform Reliable Broadcast specification
4. If  $p_3$  does not deliver  $m_5$ , then the resulting run satisfies the Uniform Reliable Broadcast specification
5. If  $p_3$  does not crash, then the resulting run satisfies the Best Effort Broadcast specification
6. The run satisfies TO(NUA, WNUTO)
7. The run satisfies the FIFO Broadcast specification
8. Let us assume  $p_3$  does not deliver  $m_5$ , then the resulting run satisfies TO(UA, WNUTO)
9. Let us assume  $p_2$  does not crash, then the resulting run satisfies TO(NUA, WNUTO)
10. If  $p_2$  does not deliver  $m_2$  and  $m_3$ , then the resulting run satisfies TO(NUA, WUTO)

- 1) **T**: VALIDITY, NO DUPLICATION AND NO CREATION ARE SATISFIED.
- 2) **F**:  $m_5$  IS NOT DELIVERED BY  $p_1$ ,  $p_2$  AND  $p_4$ .
- 3) **F**: SAME AS 2.
- 4) **T**: ALL MSGs DELIVERED (BY FAULTY OR CORRECT PROCESSES) ARE DELIVERED BY ALL CORRECT PROCESSES.
- 5) **F**: VALIDITY IS NOT SATISFIED.
- 6) **T**: NUA:  $m_5$  IS NOT DELIVERED BY CORRECT PROCESSES.  
WNUTO: IF  $p$  AND  $q$ , CORRECT PROCESSES, BOTH DELIVER  $m$  AND  $m'$ ,  $p$  DELIVERS  $m$  BEFORE  $m'$  ONLY IF  $q$  DELIVERS  $m$  BEFORE  $m'$
- 7) **T**:  $m_2 \rightarrow m_3$  SATISFIED (ONLY CORRECT)
- 8) **T**: SAME AS 4
- 9) **F**: WNUTO NOT SATISFIED (EX.  $m_1$  AND  $m_2$ )
- 10) **T**: WUTO SATISFIED

Ex. 3: Let us consider the following execution history



Assuming that the initial value stored in the register is 0, assess the truthfulness of every statement and provide a motivation for your answer:

1. If the proposed run refers to a regular register, then 0 is a valid value for read<sub>1</sub>
2. If the proposed run refers to a regular register, then 0 is not a valid value for read<sub>2</sub>
3. If the proposed run refers to a regular register, then a is a valid value for read<sub>4</sub>
4. If the proposed run refers to a regular register, then x can be returned only by read<sub>2</sub>
5. If the proposed run refers to a regular register, then read<sub>5</sub> may return only b
6. If the proposed run refers to an atomic register, then read<sub>3</sub> and read<sub>2</sub> must return the same value
7. If the proposed run refers to an atomic register, then read<sub>5</sub> returns b if and only if read<sub>3</sub> returns b
8. If the proposed run refers to an atomic register, then read<sub>6</sub> and read<sub>5</sub> must always return the same value
9. If the proposed run refers to an atomic register and read<sub>3</sub> returns b, then read<sub>6</sub> necessarily returned b
10. If the proposed run refers to an atomic register and read<sub>2</sub> returned a then read<sub>4</sub> can return only the value a

- 1) F: R<sub>1</sub>() : x, a, b
- 2) F: R<sub>2</sub>() : 0, x, a
- 3) T: R<sub>4</sub>() : x, a
- 4) F: x CAN BE RETURNED BY R<sub>1</sub>, R<sub>2</sub>, R<sub>3</sub>, R<sub>4</sub>
- 5) F: R<sub>5</sub>() : a, b
- 6) F: THEY ARE NOT CONCURRENT
- 7) F: R<sub>5</sub>() : b ALSO IF R<sub>3</sub>() : a
- 8) F: THEY ARE CONCURRENT AND RETURN SAME VALUES, BUT NOT ALWAYS. IT CAN BE R<sub>6</sub>() : b , R<sub>5</sub>() : a
- 9) T
- 10) T



Ex. 4: Let us consider the following algorithm implementing a publish-subscribe primitive

#### Algorithm 1: Publisher code

```

1 Init
2 brokers[] ← get.brokers(); /* array of length |T| where in the i-th entry the
   publisher stores the set of brokers managing topic  $t_i$  */
3 upon event publish( $e, t_i$ )
4   foreach  $b_j \in \text{brokers}[t_i]$  do
5     trigger pp2pSend(PUBLISH,  $e, t_i$ ) to  $b_j$ ;
6 end

```

#### Algorithm 2: Subscriber code

```

1 Init
2 brokers[] ← get.brokers(); /* array of length |T| where in the i-th entry the
   subscriber stores the set of brokers managing topic  $t_i$  */
3 my_subscription ← get.my_initial_subscription(); /* set of topics */
4 notified[] ← [0]^{|T|}; /* array of length |T| where in the i-th entry the
   subscriber stores the set of notified events for topic  $t_i$  */
5 upon event subscribe( $t_i$ )
6   my_subscription ← my_subscription  $\cup \{t_i\}$ ;
7   foreach  $b_j \in \text{brokers}[t_i]$  do
8     trigger pp2pSend(SUBSCRIBE,  $s_i, t_i$ ) to  $b_j$ ;
9 end
10 upon event unsubscribe( $t_i$ )
11   my_subscription ← my_subscription  $\setminus \{t_i\}$ ;
12   foreach  $b_j \in \text{brokers}[t_i]$  do
13     trigger pp2pSend(UNSUBSCRIBE,  $s_i, t_i$ ) to  $b_j$ ;
14 end
15 upon event pp2pDeliver(NOTIFY,  $e, t_i$ )
16   if  $t_i \in \text{my\_subscription}$  AND  $e_j \notin \text{notified}[t_i]$  then
17     trigger NOTIFY( $e, t_i$ );
18     notified ← notified  $\cup \{e_j, t_i\}$ ;
19 end

```

#### Algorithm 3: Broker code

```

1 Init
2 subscription[] ← [0]^{|T|}; /* array of length |T| where in the i-th entry the
   broker stores the set of subscribers for topic  $t_i$  */
3 topic_to_manage ← get.topic_to_manage(); /* set of topics that the current
   broker must manage */
4 upon event pp2pDeliver(PUBLISH,  $e, t_i$ )
5   if  $t_i \in \text{topic\_to\_manage}$  then
6     foreach  $s_j \in \text{subscription}[t_i]$  do
7       trigger pp2pSend(NOTIFY,  $e, t_i$ ) to  $s_j$ ;
8     end
9 end
10 upon event pp2pDeliver(SUBSCRIBE,  $s_i, t_i$ )
11   subscription[t_i] ← subscription[t_i]  $\cup \{s_i\}$ ;
12 upon event pp2pDeliver(UNSUBSCRIBE,  $s_i, t_i$ )
13   subscription[t_i] ← subscription[t_i]  $\setminus \{s_i\}$ ;

```

Let us assume that

1. the underlying communication system is synchronous and that every message sent over perfect point to point link takes at most  $\delta$  time to be delivered.
2. links guarantee the FIFO property
3. for each topic  $t_i$ , there are at least 3 brokers managing it

Asses the truthfulness of every statement and provide a motivation for your answer:

1. If subscriptions are not modified and there is at most one failure in the system, then every event  $e$  published on a topic  $t_i$  will be eventually notified to all subscribers subscribed to it
2. If subscriptions are not modified and there are no failures in the system, then any pair of subscribers  $s_i$  and  $s_j$  both subscribed to the same topic  $t_i$  will deliver the same set of events published on  $t_i$
3. If subscriptions are not modified and there is at most one failure in the system, if a subscriber  $s_i$  notifies an event  $e$  before an event  $e'$  (both published on a topic  $t_i$ ), then any other subscriber having  $t_i$  in its subscriptions will notify  $e$  before  $e'$
4. In absence of failures, every event published on a topic  $t_i$  at time  $\tau$  will be eventually notified to every subscriber having  $t_i$  in its subscription at time  $\tau$
5. In absence of failures, if a subscribers  $s_i$  unsubscribes a topic  $t_i$  at time  $\tau$ , it will not notify any event for topic  $t_i$  after time  $\tau$

- 1) **F**: IF A SUBSCRIBER FAILS, THEN HE WILL NOT RECEIVE ANY NEW EVENT.
- 2) **T**: SINCE FIFO PROPERTY, PERFECT P2P LINKS AND NO FAILURES ARE GUARANTEED,  $S_i$  AND  $S_j$  WILL DELIVER THE SAME SET IN THE SAME ORDER.
- 3) **T**: SINCE FIFO PROPERTY IS GUARANTEED, IF  $S_i$  NOTIFIES  $e$  BEFORE  $e'$  THEN EVERY  $S$  OF  $T_i$  NOTIFIES  $e$  BEFORE  $e'$ , EVEN WITH A FAILURE.
- 4) **F**: TIMES CAN BE DIFFERENT.
- 5) **F**: IF THERE IS A DELAY IN UNSUBSCRIBE,  $S_i$  COULD RECENE SOME EVENT PUBLISHED JUST BEFORE.

Ex. 5:

Let us consider a distributed system composed of  $n$  processes  $p_1, p_2, \dots, p_n$  partitioned in two disjoint sets acting as clients and replicas. Each process is identified by a unique integer identifier and can communicate with other processes by exchanging messages using perfect point-to-point links. Processes collaborate to implement a replication protocol emulating a linearizable append-only log.

The append-only log is characterized by the following operations:

- **Append( $v$ )**: appends the value  $v$  to the queue of the log
- **Get**: returns the content of the log, e.g. the list of values in the order in which they have been appended.

Assuming that:

1. processes may fail by crash
2. Every pair of processes can communicate exchanging messages over perfect point-to-point links
3. All processes have access to a perfect failure detector
4. Replicas have access to a uniform reliable broadcast primitive

Write the pseudocode of the distributed algorithm implementing the append-only log.