

Dependable Distributed Systems
Master of Science in Engineering in Computer Science


AA 2024/2025

Lecture 10 – Exercises

Ex 1: Let channel A and channel B be two different types of point-to-point channels satisfying the following properties:

- channel A: if a correct process p_i sends a message m to a correct process p_j at time t , then m is delivered by p_j by time $t + \delta$.
- channel B: if a correct process p_i sends a message m to a correct process p_j at time t , then m is delivered by p_j with probability p_{cons} ($p_{\text{cons}} < 1$).

Let us consider the following systems composed by 4 processes p_1, p_2, p_3 and p_4 connected through channels A and channels B.



Assuming that each process p_i is aware of the type of channel connecting it to any other process, answer to the following questions:

1. is it possible to design an algorithm implementing a perfect failure detector in system 2 when only processes having an outgoing channel of type B can fail by crash?
2. is it possible to design an algorithm implementing a perfect failure detector in system 2 if any process can fail by crash?
3. is it possible to design an algorithm implementing a perfect failure detector in system 3?

For each point, if an algorithm exists write its pseudo-code, otherwise show the impossibility.

Ex 2: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$ that communicate by exchanging messages on top of a line topology, where p_1 and p_n are respectively the first and the last process of the network.

Initially, each process knows only its left neighbors and its right neighbor (if they exist) and stores the respective identifiers in two local variables LEFT and RIGHT. Processes may fail by crashing, but they are equipped with a perfect oracle that notifies at each process the new neighbor (when one of the two fails) through the following primitives:


- **Left_neighbour(p_x):** process p_x is the new left neighbor of p_i
- **Right_neighbour(p_x):** process p_x is the new right neighbor of p_i

Both the events may return a NULL value in case p_i becomes the first or the last process of the line.

Each process can communicate only with its neighbors.

Write the pseudo-code of an algorithm implementing a Leader Election primitive assuming that channels connecting two neighbor processes are perfect.

Ex 3: Consider the partial execution depicted in the Figure



Answer to the following questions:

1. Complete the execution in order to have a run satisfying Uniform Reliable Broadcast.
2. Complete the execution in order to have a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast.
3. Complete the execution in order to have a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast.

NOTE: In order to solve the exercise, you can add broadcast, deliver and crash events but you cannot remove anything from the run.

Ex 4: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$ identified through unique integer identifiers. Processes may communicate using perfect point-to-point links. Links are available between any pair of processes. Processes may fail by crash and each process has access to a perfect failure detector.

Modify the algorithms implementing a distributed mutual exclusion abstraction discussed during the lectures to allow them to tolerate crash failures.

Leader based Algorithm

PROCESS CODE

Init

```
state = idle
coordinator = getCoordinatorId()
correct = {p1, p2, ... pn}
transition = false
```

upon event request()

```
state = waiting
trigger pp2pSend(REQ, i) to coordinator
```

upon event pp2pDeliver(GRANT_CS)

```
state = CS
trigger ok()
```

upon event release()

```
state = idle
trigger pp2pSend(REL, i) to coordinator
```

upon event crash(pj)

```
correct = correct \ {pj}
if coordinator == pj
    coordinator = getCoordinatorId(correct)
    if coordinator == self
        % manage transition
        transition = true
        pending =  $\emptyset$ 
        current_in_CS =  $\perp$ 
        ACK = NACK =  $\emptyset$ 
        for every pk  $\in$  correct
            trigger pp2pSend(CHECK_CS) to pk
    else
        if state == waiting
            trigger pp2pSend(REQ, i) to coordinator
```

upon event pp2pDeliver(CHECK_CS) from coordinator

```
if state == CS
    trigger pp2pSend(ACK, self) to coordinator
else
    trigger pp2pSend(NACK, self) to coordinator
```

COORDINATOR CODE

Init

```
pending =  $\emptyset$ 
```

current_in_CS= \perp
correct = {p1, p2, ... pn}
transition = false

upon event pp2pDeliver(REQ, j) from p_j
 pending = pending \cup {p_j}

when pending $\neq \emptyset$ and current_in_CS= \perp and not transition
 candidate=select_process (pending)
 pending = pending \setminus candidate
 current_in_CS = candidate
 trigger pp2pSend(GRANT_CS) to candidate

upon event pp2pDeliver(REL, j) from p_j
 if current_in_CS == j
 current_in_CS = \perp

upon event crash(p_j)
 correct = correct \setminus {p_j}
 if p_j \in pending
 pending = pending \setminus {p_j}
 if current_in_CS= p_j
 current_in_CS= \perp

upon event pp2pDeliver (ACK, j)
 current_in_CS = j
 transition= false


upon event pp2pDeliver (NACK, j)
 NACK= NACK \cup {j}

when correct \subseteq NACK
 transition = false

Ex 1: Let channel A and channel B be two different types of point-to-point channels satisfying the following properties:

- channel A: if a correct process p_i sends a message m to a correct process p_j at time t , then m is delivered by p_j by time $t + \delta$
- channel B: if a correct process p_i sends a message m to a correct process p_j at time t , then m is delivered by p_j with probability p_{cons} ($p_{\text{cons}} < 1$).

Let us consider the following systems composed by 4 processes p_1, p_2, p_3 and p_4 connected through channels A and channels B.



Assuming that each process p_i is aware of the type of channel connecting it to any other process, answer to the following questions:

- is it possible to design an algorithm implementing a perfect failure detector in system 2 when only processes having an outgoing channel of type B can fail by crash?
- is it possible to design an algorithm implementing a perfect failure detector in system 2 if any process can fail by crash?
- is it possible to design an algorithm implementing a perfect failure detector in system 3?

For each point, if an algorithm exists write its pseudo-code, otherwise show the impossibility.

1) YES BECAUSE AT LEAST ONE PROCESS RECEIVE ALWAYS THE RIGHT INFORMATION.

INIT

```

ALIVE =  $\Pi$ 
DETECTED =  $\emptyset$ 
TIMER =  $\Delta$ 
PHASE = 1
STARTTIMER (TIMER)
  
```

```

UPON EVENT < TIMEOUT AND PHASE == 1 > DO
  FORALL  $p \in \Pi$  DO
    TRIGGER SEND (LIST.ALIVE) TO  $p$ 
  PHASE == 2
  STARTTIMER (TIMER)
  
```

```

UPON EVENT < TIMEOUT AND PHASE == 2 > DO
  FORALL  $p \in \Pi$  DO
    IF ( $p \notin \text{ALIVE} \wedge p \notin \text{DETECTED}$ ) THEN
      DETECTED = DETECTED  $\cup \{p\}$ 
      TRIGGER < CRASH |  $p$  >
      TRIGGER < SEND | [HBRQ] >
  ALIVE =  $\emptyset$ 
  PHASE = 1
  STARTTIMER (TIMER)
  
```

```

UPON EVENT DELIVER (HBRQ) FROM  $q$ 
  TRIGGER SEND (HBRP) TO  $q$ 
  
```

```

UPON EVENT DELIVER (HBRP) FROM  $q$ 
  ALIVE = ALIVE  $\cup \{q\}$ 
  
```

```

UPON EVENT DELIVER (LIST, ALIVE.LIST)
  ALIVE = ALIVE  $\cup \{\text{ALIVE.LIST}\}$ 
  
```

2) NO, BECAUSE IN THIS CASE WE CANNOT HAVE A PROCESS THAT HAS THE COMPLETE VIEW OF THE ALIVE PROCESSES. SO WE CAN DO IT UNTIL THE PROCESS p_i REMAIN CORRECT.

3) NO, BECAUSE IN THIS CASE THE SYSTEM IS COMPOSED ONLY BY CHANNEL B THAT IS FAIR LOSS LNK, AND WE CANNOT IMPLEMENT PFD OVER FAIR LOSS

Ex 2: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$ that communicate by exchanging messages on top of a line topology, where p_1 and p_n are respectively the first and the last process of the network.


Initially, each process knows only its left neighbors and its right neighbor (if they exist) and stores the respective identifiers in two local variables LEFT and RIGHT. Processes may fail by crashing, but they are equipped with a perfect oracle that notifies at each process the new neighbor (when one of the two fails) through the following primitives:

- **Left_neighbour(p_x):** process p_x is the new left neighbor of p_i
- **Right_neighbour(p_x):** process p_x is the new right neighbor of p_i

Both the events may return a NULL value in case p_i becomes the first or the last process of the line.

Each process can communicate only with its neighbors.

Write the pseudo-code of an algorithm implementing a Leader Election primitive assuming that channels connecting two neighbor processes are perfect.



UPON EVENT $\langle le, INIT \rangle$ DO

LEFT: GET.LEFT()

RIGHT: GET.RIGHT()

LEADER = p_i

UPON EVENT LEFT.NEIGHBOUR(p_x)

LEFT = p_x

IF LEFT = NULL

LEADER = MY.ID

TRIGGER LEADER(LEADER)

TRIGGER PP2PSEND(NEW.LEADER, LEADER) TO RIGHT

UPON EVENT RIGHT.NEIGHBOUR(p_x)

RIGHT = p_x


UPON EVENT PP2PDELIVER(NEW.LEADER, p_x) FROM LEFT

LEADER = p_x

TRIGGER LEADER(LEADER)

TRIGGER PP2PSEND(NEW.LEADER, LEADER) TO RIGHT

Ex 3: Consider the partial execution depicted in the Figure



Answer to the following questions:

1. Complete the execution in order to have a run satisfying Uniform Reliable Broadcast.
2. Complete the execution in order to have a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast.
3. Complete the execution in order to have a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast.

NOTE: In order to solve the exercise, you can add broadcast, deliver and crash events but you cannot remove anything from the run.

1) URB

The diagram shows the completion of the execution for URB. All messages m1 through m5 are delivered to all correct processes (P1, P2, P4). P3 has crashed. A legend indicates that a double arrow represents a broadcast and a single arrow represents a deliver.

ALL CORRECT PROCESSES MUST DELIVER EVERY MESSAGE THAT ANOTHER PROEISS (CORRECT OR FAILED) HAS DELIVERED.

2) RB BUT NO URB

The diagram shows the completion of the execution for RB. All messages m1 through m5 are delivered to all correct processes (P1, P2, P4). P3 has crashed. A legend indicates that a double arrow represents a broadcast and a single arrow represents a deliver.

P₁, P₂ AND P₄ DON'T DELIVER m₆ BECAUSE IT'S ONLY DELIVERED BY P₃, WHICH IS A CRASHED PROCESS.

3) BEB BUT NO RB

The diagram shows the completion of the execution for BEB. All messages m1 through m5 are delivered to all correct processes (P1, P2, P4). P3 has crashed. A legend indicates that a double arrow represents a broadcast and a single arrow represents a deliver.

m₅ NO AGREEMENT

Ex 4: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$ identified through unique integer identifiers. Processes may communicate using perfect point-to-point links. Links are available between any pair of processes. Processes may fail by crash and each process has access to a perfect failure detector.


Modify the algorithms implementing a distributed mutual exclusion abstraction discussed during the lectures to allow them to tolerate crash failures.

INIT

```

ID ← IDENTIFIER
PENDING = ∅
REPLIES = ∅
CORRECT = Π
STATE = NCS
LAST.REQ = 0

```



UPON EVENT < ME, REQUEST > DO

```

STATE = REQUESTING
LAST.REQ = ID
FORALL q ∈ Π DO
    TRIGGER < PP2PL, SEND | REQUEST, ID >

```

UPON EVENT < PP2PL, DELIVER | REQUEST, q >

```

IF (STATE == CS) OR (STATE == REQUESTING ∧ LAST.REQ ≠ ID)
    PENDING = PENDING ∪ {q}
ELSE
    TRIGGER < PP2PL, SEND | [REPLAY] >

```

UPON EVENT < PP2PL, DELIVER | [REPLAY], q >

```

REPLIES = REPLIES ∪ {q}

```

UPON EVENT < P, CRASH | p >

```

CORRECT = CORRECT \ {p}

```

WHEN REPLIES == CORRECT DO

```

STATE = CS
TRIGGER < ME, CS >
FORALL q ∈ PENDING
    TRIGGER < PP2PL, SEND | [REPLAY] > TO q
STATE = NCS
REPLIES = ∅

```