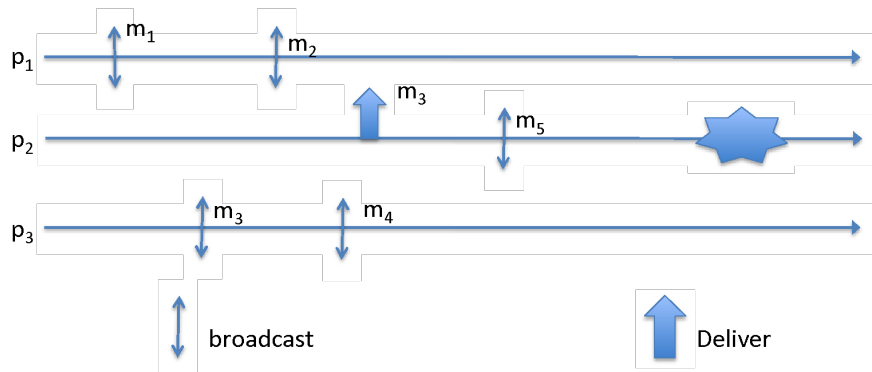


Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2024/2024

Lecture 13 – Exercises
October 20th, 2024

Ex 1: Consider the partial execution depicted in the following figure:



1. Complete the execution in order to obtain a run satisfying *Best Effort Broadcast* but *not Reliable Broadcast*.
2. Complete the execution in order to obtain a run satisfying *Regular Reliable Broadcast* but *not Uniform Reliable Broadcast*.
3. Complete the execution in order to obtain a run satisfying *Uniform Reliable Broadcast*.

Ex 2: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$. Each process is connected to all the others through fair-loss point-to-point links and has access to a perfect failure detector.

Write the pseudo-code of an algorithm implementing a Uniform Reliable Broadcast primitive.

Given the system model described here, additionally answer to the following questions:

1. Is it possible to provide a quiescent implementation of the Uniform Reliable Broadcast primitive (a quiescent implementation is one where all processes eventually stop sending messages)?
2. Is it possible to provide an implementation that uses only data structure with finite size?

Ex 3: Consider a distributed system composed by N servers $\{s_1, s_2, \dots, s_n\}$ and M clients $\{c_1, c_2, \dots, c_m\}$.

Each client c_i runs its algorithm and it can request to servers the execution of a particular task T_i . A Server will execute the task T_i and, after that, a notification will be sent to c_i that T_i has been completed.

The Figure shows the code executed by a generic client c_i .

Operation executeTask (T_i) 1. For each $s_i \in \{s_1, s_2, \dots, s_n\}$ 2. pp2psend (TASK_REQ, T_i , c_i) to s_i ;	Upon pp2pdeliver (TASK_COMPLETED, T_i) from s_j 1. trigger completedTask (T_i);
--	--


Write the pseudo-code of an algorithm, executed by servers, able to allocate tasks assuming that:

- Once clients ask for a task execution, they remain blocked until the task is not terminated.
- Any two clients c_i and c_j can concurrently require the execution of two different tasks T_i and T_j ;
- Each task is univocally identified by the pair (T_i, c_i) ;
- Each server can manage at most one task at every time;
- At most $N-1$ servers can crash;
- If a server crashes while executing a task, such task needs to be re-allocated and re-processed by a different server;
- Servers have access to a uniform consensus primitive;
- Servers have access to a perfect failure detector P ;
- Servers communicate through a uniform reliable broadcast primitive.

Ex 4: Consider a distributed system formed by n processes p_1, p_2, \dots, p_n connected along a ring i.e., a process p_i is initially connected to a process $p_{(i+1) \bmod n}$ through a unidirectional perfect point-to-point link.


Write the pseudo-code of a distributed algorithm implementing a consensus primitive.

Ex 1: Consider the partial execution depicted in the following figure:




1. Complete the execution in order to obtain a run satisfying *Best Effort Broadcast* but *not Reliable Broadcast*.
2. Complete the execution in order to obtain a run satisfying *Regular Reliable Broadcast* but *not Uniform Reliable Broadcast*.
3. Complete the execution in order to obtain a run satisfying *Uniform Reliable Broadcast*.


1) BEB BUT NOT RB



2) RRB BUT NOT URB



3) URB



Ex 2: Consider a distributed system composed by n processes $\{p_1, p_2, \dots, p_n\}$. Each process is connected to all the others through fair-loss point-to-point links and has access to a perfect failure detector.


Write the pseudo-code of an algorithm implementing a Uniform Reliable Broadcast primitive.

Given the system model described here, additionally answer to the following questions:

1. Is it possible to provide a quiescent implementation of the Uniform Reliable Broadcast primitive (a quiescent implementation is one where all processes eventually stop sending messages)?
2. Is it possible to provide an implementation that uses only data structure with finite size?

```

UPON EVENT <URB, INIT> DO
    PENDING =  $\emptyset$ 
    DELIVERED =  $\emptyset$ 
    CORRECT =  $\Pi$ 
    FORALL  $m$  DO  $ACK[m] = \emptyset$ 
    STARTTIMER ( $\Delta$ )
  
```



```

UPON EVENT <URB, BROADCAST | m> DO
    PENDING = PENDING  $\cup \{m\}$ 
    FORALL  $p_i$  IN  $\Pi$  DO
        TRIGGER <FL, FL.SEND | (MSG, m,  $p_i$ )> TO  $p_i$ 
  
```

```

UPON EVENT <FL, FL.DELIVER | (MSG, m,  $p_i$ )> FROM  $p_i$  DO
     $ACK[m] = ACK[m] \cup \{p_i\}$ 
    FORALL  $p_i$  IN  $\Pi$  DO
        TRIGGER <FL, FL.SEND | (MSG, m,  $p_i$ )> TO  $p_i$ 
  
```

```

UPON EVENT <PFD, CRASH |  $p_j$ > DO
    CORRECT = CORRECT  $\setminus \{p_j\}$ 
  
```

```

UPON EXIST  $m$  IN PENDING SUCH THAT CORRECT  $\subseteq ACK[m] \wedge m$  NOT IN DELIVERED DO
    DELIVERED = DELIVERED  $\cup \{m\}$ 
    TRIGGER <URB, DELIVER | m>
  
```

```

UPON <TIMEOUT>
    FORALL  $m$  IN PENDING DO
        FORALL  $p_i$  IN  $\Pi$  DO
            TRIGGER <FL, FL.SEND | (MSG, m,  $p_i$ )> TO  $p_i$ 
    STARTTIMER ( $\Delta$ )
  
```

Ex 3: Consider a distributed system composed by N servers $\{s_1, s_2, \dots, s_n\}$ and M clients $\{c_1, c_2, \dots, c_m\}$.

Each client c_i runs its algorithm and it can request to servers the execution of a particular task T_i . A Server will execute the task T_i and, after that, a notification will be sent to c_i that T_i has been completed.

The Figure shows the code executed by a generic client c_i .

Operation executeTask (T_i)	Upon pp2pldeliver (TASK_COMPLETED, T_i) from s_j
<ol style="list-style-type: none"> For each $s_i \in \{s_1, s_2, \dots, s_n\}$ pp2psend (TASK_REQ, T_i, c_i) to s_i; 	<ol style="list-style-type: none"> trigger completedTask (T_i);

Write the pseudo-code of an algorithm, executed by servers, able to allocate tasks assuming that:

- Once clients ask for a task execution, they remain blocked until the task is not terminated.
- Any two clients c_i and c_j can concurrently require the execution of two different tasks T_i and T_j ;
- Each task is univocally identified by the pair (T_i, c_i) ;
- Each server can manage at most one task at every time;
- At most $N-1$ servers can crash;
- If a server crashes while executing a task, such task needs to be re-allocated and re-processed by a different server;
- Servers have access to a uniform consensus primitive;
- Servers have access to a perfect failure detector P ;
- Servers communicate through a uniform reliable broadcast primitive.

CLIENT:

OPERATION EXECUTE TASK (T_i)

FOR EACH $s_i \in \{s_1, \dots, s_n\}$

PP2PL.SEND (TASK.REQ, T_i, c_i) TO s_i

UPON PP2PL.DELIVER (TASK_COMPLETED, T_i) FROM s_j

TRIGGER COMPLETEDTASK (T_i)

SERVER:

INIT

PENDING = \emptyset

BUSY = FALSE

WARENT.ALLOCATION = \emptyset

UPON EVENT \langle PP2PL.DELIVER (TASK.REQ, T_i, c_i) \rangle FROM c_i

IF $\langle T_i, c_i \rangle$ NOT IN WARENT.ALLOCATION

PENDING = PENDING $\cup \{ (T_i, c_i) \}$

WHEN PENDING $\neq \emptyset$

IF !BUSY

CANDIDATE = SELECT.FROM (PENDING)

TRIGGER PROPOSE (\langle ID, CANDIDATE \rangle)

ELSE

TRIGGER PROPOSE (\langle ID, NULL \rangle)

UPON EVENT CRASH (s)

IF THERE EXISTS \langle ID, TASK \rangle IN WARENT.ALLOCATION SUCH THAT $s ==$ ID

WARENT.ALLOCATION = WARENT.ALLOCATION $\setminus \{$ ID, TASK $\}$

PENDING = PENDING $\cup \{ (T_i, c_i) \}$

UPON EVENT DECIDE (ID, TASK)

WARENT.ALLOCATION = WARENT.ALLOCATION $\cup \{$ ID, TASK $\}$

IF ID == MY.ID AND TASK \neq NULL

BUSY = TRUE

EXECUTE (T_i)

BUSY = FALSE

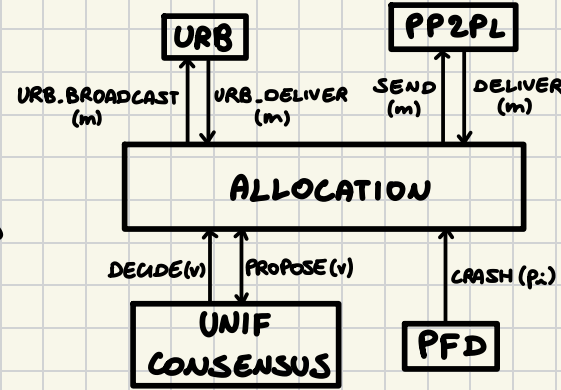
TRIGGER \langle PP2PL.SEND (TASK.COMPLETED, T_i) \rangle TO c_i

TRIGGER \langle URB.BROADCAST (TASK.COMPLETED, T_i, c_i)

UPON EVENT URB.DELIVER (TASK.COMPLETED, T_i, c_i)

WARENT.ALLOCATION = WARENT.ALLOCATION $\setminus \{ T_i, c_i \}$



PENDING = PENDING $\setminus \{ T_i, c_i \}$



Ex 4: Consider a distributed system formed by n processes p_1, p_2, \dots, p_n connected along a ring i.e., a process p_i is initially connected to a process $p_{(i+1) \bmod n}$ through a unidirectional perfect point-to-point link.

Write the pseudo-code of a distributed algorithm implementing a consensus primitive.

$P_{(i+1) \bmod n}$



UPON EVENT $\langle C, \text{INIT} \rangle$ DO
 NEXT = $P_{(i+1) \bmod n}$
 DECISION = \perp
 PROPOSAL.SET = \emptyset
 ALIVE = \top

UPON EVENT $\langle C, \text{PROPOSE } |v\rangle$ DO
 PROPOSAL = v
 PROPOSAL.SET = PROPOSAL.SET $\cup \{v\}$
 TRIGGER $\langle \text{PP2PL}, \text{PP2PL.SEND} | (\text{PROPOSAL}, \text{PROPOSAL.SET}, p_i) \rangle$ TO NEXT

UPON EVENT $\langle \text{PP2PL}, \text{PP2PL.DELIVER} | (\text{PROPOSAL}, \text{PROPOSAL.SET}, p_i) \rangle$ FROM p_j
 IF $p_i = p_j$ THEN
 IF ALIVE \leq PROPOSAL.SET \wedge DECISION = \perp DO
 DECISION = $\min(\text{PROPOSAL.SET})$
 TRIGGER $\langle C, \text{DECIDE } | \text{DECISION} \rangle$
 ELSE
 PROPOSAL.SET = $p_s \cup p_s$
 TRIGGER $\langle \text{PP2PL}, \text{PP2PL.SEND} | (\text{PROPOSAL}, \text{PROPOSAL.SET}, p_i) \rangle$ TO NEXT

